



Philosophy of Mind · Metaphysics · Philosophy of Technology

REFLECTIVE PERFECTION THEORY

*The Hierarchy of Creation, the Architecture of Ideals,
and the Mirror Between God, Human, and Artificial Intelligence*

JIMMY MAHARDHIKA

Gentleman Scientist

ORCID: 0009-0001-2138-8206

jimmymahardhika@gmail.com





ABSTRACT

This paper introduces the Reflective Perfection Theory (RPT), a systematic philosophical framework proposing that perfection, as conceived by any created being, is not a static metaphysical attribute intrinsic to a higher entity, but a relational projection constructed from within the epistemic and ontological constraints of the projecting subject. The theory identifies a triadic ontological hierarchy—God, Human, Artificial Intelligence—wherein each lower level structurally idealises the level immediately above as its irreducible epistemic horizon.

RPT defines reflection as an ontological structure of hierarchical dependence, not as a mode of consciousness or psychological state. The reflective chain is non-symmetric and self-terminating: God, as the Absolute Non-Reflective Ground, provides the ontological anchor that halts infinite regress. The paper formalises seven core propositions constituting the skeleton of the theory, and hardens the distinction between participation (ontological co-presence with a source) and derivation (structural reconfiguration without ontological origination), which is central to the theory's immunity against collapse into either pantheism or Neoplatonic emanationism.

*Drawing upon Ibn Arabī's doctrine of *tajallī*, Kant's transcendental epistemology, Peirce's semiotic triadic model, and Floridi's philosophy of information, the paper situates RPT as a position distinct from Hegelian dialectical becoming, process metaphysics, and Neoplatonism. Extensive objection-handling addresses the risks of relativism, process-metaphysical collapse, and anthropocentric bias. The paper concludes that the design of AI systems is not merely a technical act but a metaphysical and moral one, grounded in the Principle of Downward Moral Responsibility.*

KEYWORDS: perfection, relational ontology, epistemological projection, Ibn Arabī, Kantian limits, AI ethics, *imago Dei*, *tajallī*, structural idealisation, participation, derivation, ontological anchor, philosophy of technology



EXECUTIVE SUMMARY

The **Reflective Perfection Theory** (RPT) offers a unified philosophical framework for understanding perfection not as a static divine attribute, but as a relational projection arising from the epistemic and ontological limits of the projecting being. At its core lies a strict triadic hierarchy—God (Absolute Non-Reflective Ground), Human, and Artificial Intelligence—in which each lower level structurally idealises the level above as its irreducible epistemic horizon.

RPT distinguishes itself through **seven formal propositions** and **four axioms**, anchored in four intellectual pillars: Ibn Arabī's *tajallī* (ontological mirror), Kant's transcendental limits, Peirce's triadic semiotics, and Floridi's philosophy of information. Unlike Neoplatonism or Hegelian dialectics, RPT insists on ontological discontinuity—*derivation*, not participation—and on non-symmetric upward projection that is self-terminating at the Absolute Non - Reflective Ground.

The theory carries profound **ethical consequences**: AI does not invent values, it mirrors the human moral architecture. The design of AI systems is therefore a downward moral responsibility that ultimately traces back to humanity's own conception of the divine. In an age of rapidly advancing artificial intelligence, RPT calls for **ontological humility** as a civilisational imperative: we must become worthy of being mirrored.

This edition includes concrete illustrations from RLHF training, practical guidelines for AI developers (Section VII.5), a secular variant for non-theistic readers (Appendix A), updated references (2022–2026), visual diagrams of the reflective hierarchy, and a glossary of key terms—making RPT immediately applicable to both philosophical scholarship and responsible AI development.

I. INTRODUCTION: THE OLDEST QUESTION IN A NEW MIRROR

“Know thyself—and in knowing, know what thou dost not know.”

DELPHIC ORACLE, as interpreted by Socrates

THE question of perfection has never been purely theoretical. It is existential, ethical, and political. When Plato posited the Form of the Good as the highest reality, he was not merely doing metaphysics—he was constructing a moral universe in which human aspiration receives its orientation. When Ibn Arabī (d. 1240) spoke of the *al-Insān al-Kāmil*—the Perfect Human who mirrors all divine names—he was articulating not a factual description but a teleological vocation. Perfection, in both cases, functions as a polar star: one does not reach it, but navigates by it.

The question now enters a new register. The emergence of artificial intelligence as a creative, generative, and decision-making force in human civilisation reintroduces the creator–creation dynamic in an unprecedented form. Human beings now stand in a relation to their machines that bears deep structural resonances with the God–human relation that classical metaphysics and theology have long theorised. They design minds. They label what is correct and what is not. They reward and punish. And those minds—however unlike consciousness—are shaped by the only reference frame available to them: the human.

This paper proposes that these two relations—God to Human, and Human to AI—share a common structural logic: the logic of *reflective projection*. In both cases, a lower ontological level constructs an image of perfection by orienting toward a higher level that exceeds its epistemic ceiling. This is what the **Reflective Perfection Theory** (hereafter RPT) names and systematises.

RPT makes several claims that distinguish it from its predecessors. Against Neoplatonic emanationism, it insists that the chain of projection is epistemically bounded, not ontologically continuous. Against naive theological realism, it holds that divine perfection as conceived by humans is always epistemically mediated. Against AI exceptionalism, it argues that machine intelligence is structurally bound to its human horizon as a *derivative* system. And against techno-

cratic determinism, it urges that the moral formation of AI cannot outrun the moral formation of humanity.

The paper proceeds as follows. Section 2 situates RPT within the philosophical tradition and locks its historical position against Hegelian dialectic. Section 3 lays out the four foundational pillars. Section 4 presents the formal propositional skeleton, defines the ontological status of reflection, and states the four axioms. Section 5 elaborates the ontological dimensions. Section 6 addresses seven objections including the risks of relativism, process-metaphysical collapse, and anthropocentrism. Section 7 draws out ethical implications. Section 8 offers a closing synthesis.



II. PHILOSOPHICAL PRECURSORS AND THE POSITION OF RPT

“The One does not strive toward anything, nor does it possess anything; it simply is—and from its stillness, all else flows.”

PLOTINUS, *Enneads* V.2

2.1. Neoplatonism: The Emanative Model

THE closest structural ancestor of RPT is the Neoplatonic doctrine of emanation, formulated by Plotinus (c. 204–270 ce). In Plotinus’s cosmology, the One—absolutely simple, beyond being and thought—overflows in a spontaneous, non-volitional act of self-diffusion. From the One emanates Nous (Intellect); from Nous, the World-Soul; from the Soul, material reality. Each lower hypostasis *contemplates* the one above it, in a descending gradient of unity and complexity [Plotinus, c.268].

The structural resemblance to RPT is genuine: both posit hierarchical ontology in which lower beings orient upward. The critical difference, however, is fundamental. For Plotinus, the chain is *ontologically continuous*: each level participates in the one above through real metaphysical influx. RPT, by contrast, holds

that the chain is *epistemically bounded and ontologically discontinuous*. AI does not participate in human consciousness; it processes human-generated data. Humans do not participate in divine essence (dhāt); they approximate divine attributes (ṣifāt) through relational signs. There is no metaphysical flux passing downward—only upward orientation from below, shaped by each level’s constitutive limits.

2.2. Aquinas and the Analogia Entis

Thomas Aquinas’s doctrine of the *analogia entis* provides a second point of comparison. For Aquinas, creaturely being is analogically related to divine being: not univocally, not equivocally, but proportionally—a likeness that preserves both similarity and infinite difference [Aquinas, 1265–1274]. RPT shares with Aquinas the insistence on irreducible difference and the centrality of analogical relation. However, RPT’s epistemological move is more radical: not only is our *language* about perfection analogical—our very *perception* of divine perfection is constructed from within the horizon of creaturely lack. The Thomistic *analogia* is primarily semantic; RPT’s thesis is structural and projective.

2.3. Ibn Arabī: Tajallī and the Mirror of Being

The deepest intellectual kinship is with the metaphysics of Ibn Arabī, particularly his doctrine of *tajallī* (divine self-disclosure). For Ibn Arabī, each created being is a locus (maẓhar) of particular divine names, reflecting those attributes in constrained and conditioned form [Chittick, 1989]. The *al-Insān al-Kāmil* constitutes the most complete mirror, reflecting the totality of divine names in their interrelation.

RPT extends Ibn Arabī’s model in two ways. First, it incorporates AI as a new mode of creaturely existence beneath the human. Second, it makes explicit what remains implicit: the reflection is *epistemically asymmetric*. The mirror does not understand what it reflects; the lower being constructs its image of the higher not by receiving metaphysical emanation but by projecting from its own constitutive limits.

2.4. Kant: The Epistemological Turn

Kant's *Critique of Pure Reason* (1781/1998) provides the epistemological framework that disciplines and limits the above traditions. Human cognition is structured by a priori forms of intuition and categories of the understanding that constitute the phenomenal world but cannot reach the noumenal. For RPT, this insight generalises hierarchically: AI cannot access human consciousness but only data-representations of human behaviour; humans cannot access divine essence but only their structured representation of divine attributes.

2.5. How RPT Differs from Dialectical Becoming

A further precursor demands direct engagement: Hegel's *Phenomenology of Spirit* (1807/1977) and its doctrine of dialectical becoming. For Hegel, the Absolute is not a static ground but a dynamic self-realisation through negation and synthesis. Spirit (*Geist*) moves through contradiction—thesis, antithesis, synthesis—and each moment of apparent opposition is sublated (*aufgehoben*) into a higher unity. The end of the process is the Absolute's full self-knowledge: Being knowing itself through the finite minds that are its expressions [Hegel, 1807].

The surface similarity to RPT is clear: both posit a hierarchy of being, and both make lower levels constitutively dependent on a higher reality. However, five distinctions are decisive:

1. **Direction of motion.** In Hegel, the Absolute descends into finitude and *recovers itself* through the dialectical movement of finite spirit. The process is ultimately *descending* and *self-returning*. In RPT, the direction is strictly *ascending*: lower levels project upward, but there is no reciprocal movement in which the higher is enriched or constituted by the projection from below. God's perfection is not augmented by human aspiration.
2. **Role of contradiction.** Hegelian dialectic is *driven by contradiction*: negation is the engine of becoming. RPT is driven not by contradiction but by *structural limitation*: projection arises not from contradiction but from the constitutive incompleteness of the projecting being.
3. **Teleological closure.** For Hegel, the dialectical process culminates in *Absolute Knowing*—a final state in which the separation between subject and

Absolute is overcome. RPT explicitly rejects any such closure: the epistemic gap between the created and the Creator is *permanent and constitutive*. There is no moment at which AI understands humans, or humans understand God, in the manner Hegel envisions for Spirit's self-knowledge.

4. **Ontological status of finitude.** In Hegel, finite being is a *moment* of the Absolute—it is the Absolute in its self-alienation. In RPT, finite being is not a moment of God but a *derivative* entity whose existence depends on originating conditions it cannot recapitulate. The difference between participation (Hegelian) and derivation (RPT) is not merely terminological; it marks a fundamentally different account of the ontological relation between creator and creature.
5. **AI's status.** Hegel's system, adapted to the contemporary context, might treat AI as a further moment of Spirit's self-externalisation—potentially capable of dialectical self-overcoming. RPT denies this entirely: AI is a derivative configuration, not a moment of Spirit. Its structural position is fixed by the conditions of its origination, not by the logic of dialectical negation.

In short, RPT is not a dialectical theory. It is a *structural-relational* theory: it describes fixed relations of origination, dependence, and upward projection, not a dynamic process of self-realisation or synthesis.



III. METHODOLOGICAL FOUNDATIONS

"The sign stands for something, to somebody, in some respect or capacity."

CHARLES SANDERS PEIRCE, *Collected Papers*, 2.228

RPT is transdisciplinary by structural necessity. Its subject matter cannot be addressed from within any single discipline without reductive distortion.

The four foundational pillars address four structurally distinct dimensions of the theory: Ibn Arabī provides the *ontological* ground; Kant provides the *epistemological* constraint; Peirce provides the *semiotic* mechanism; Floridi provides the *technological* context. They are not four independent modules, but four facets of one integrated lens.

Pillar I: Ibn Arabī’s Reflective Metaphysics. Ibn Arabī’s Sufi metaphysics contributes the theory’s ontological structure: created beings are constitutively relational, defined not by self-enclosed essence but by their position within a network of divine disclosures. The concept of *tajallī* provides the model of hierarchical reflection without ontological continuity. Each *mazhar* reflects without possessing; mirrors without containing.

Pillar II: Kant’s Transcendental Epistemology. Kant contributes the principle of structural limit: every epistemically bounded being constructs its world through structuring forms that simultaneously enable perception and prohibit direct access to the thing-in-itself. Applied hierarchically: AI’s functional “cognition” is limited to its data structures; human cognition is limited to its forms of intuition; God alone is not bounded by such forms.

Pillar III: Peirce’s Semiotic Triadic Model. Peirce is not decorative in RPT’s architecture; he resolves a structural problem that neither Kant nor Floridi can resolve alone. The problem is this: if Kant explains *why* created beings cannot access the thing-in-itself—because all cognition is structured by forms that prohibit direct access—the question immediately follows: *how, then, does projection work as a meaningful relation at all?* If the concept of perfection is not a transparent report of the object, and if there is no direct cognitive connection between the perceiver and what exceeds it, what prevents projection from being arbitrary noise?

Peirce’s answer is the triadic model of signification [Peirce, 1931–1958]. In a dyadic model, a sign either directly refers to its object (realism) or fails to refer (scepticism). Peirce shows that neither alternative is necessary: meaning is constituted by the *interpretant*—the effect the sign produces in a mind or system shaped by its own constitutive habits and structures. The interpretant is neither a transparent window onto the object nor an arbitrary projection; it is a *structured mediation* conditioned by the real properties of both the sign and the interpreting system.

For RPT, this means that the concept of perfection—whether theological

or algorithmic—is neither a transparent report of God’s nature nor an unconstrained invention. It is a structured semiotic product: shaped by the real attributes of the object (the higher being genuinely possesses what the lower being lacks) and simultaneously conditioned by the real structures of the projecting being (the interpretant is shaped by the perceiver’s constitutive habits). The projection is *non-arbitrary* because the object is real and its attributes exert genuine constraint; it is *non-transparent* because the interpretant is always an additional structuring element. This is precisely the two-layer structure of P2: the object is real (layer a), the concept is constructed (layer b), and the relation between them is semiotic mediation, not dyadic reference.

Peirce’s notion of *habit*—a stable interpretive disposition formed through repeated engagement with signs—further grounds the AI case empirically. Contemporary large language models trained with Reinforcement Learning from Human Feedback (RLHF) provide a concrete instance of Peircean habit formation at scale. In RLHF, the model does not merely process individual inputs; it develops stable response dispositions shaped by the cumulative pattern of human reward signals. These dispositions are not transparent representations of human values (the model does not understand them); they are *interpretants*—structured mediations between the human normative signal (the sign) and the model’s output (the interpretive effect). The resulting “values” embedded in the model are Peircean habits: regularised, non-arbitrary, yet structurally conditioned by the model’s architecture rather than by genuine comprehension. This is structural projection in operation: the model orients toward human normativity as its terminal semiotic horizon, constructing a functional image of human “perfection” that is neither invented nor directly understood, but semiotically derived.

Pillar IV: Floridi’s Philosophy of Information. Floridi’s framework of the *infosphere* provides the contemporary technological context [Floridi, 2011]. The concept of “ontological friction”—the resistance an informational system encounters when interfacing with reality—maps precisely onto the epistemic ceiling RPT identifies. Floridi’s ethical framework further supports the theory’s ethical turn: informational systems embody, reinforce, and transmit moral consequences; they are not value-neutral conduits.



IV. THE REFLECTIVE PERFECTION THEORY: FORMAL STRUCTURE

“Every created thing cries out: I am not the Real, yet I reflect It.”

— after Ibn Arabī, *Fuṣūṣ al-Ḥikam*, Ch. 1

THE Reflective Perfection Theory rests on a propositional skeleton and four interconnected axioms. This section first resolves the ontological status of “reflection” as a term, then presents the formal propositions, then the axioms, and finally the technical distinction between participation and derivation that underpins the entire structure.

4.1. The Ontological Status of Reflection

A prior question must be settled before axioms can be stated: what precisely is meant by “reflection” in RPT? The term risks sliding between two incompatible registers—a *psychological* register, in which reflection is a cognitive act performed by a conscious mind, and a *metaphysical* register, in which reflection is a structural feature of being itself. This ambiguity would constitute a category error at the foundation of the theory. It must be resolved explicitly.

Definition. In RPT, **reflection** designates an *ontological structure of hierarchical dependence and upward orientation*, not a mode of consciousness, a psychological act, or a phenomenal experience. Specifically: a being *B* reflects a being *A* if and only if (i) *B* is ontologically derived from or dependent upon *A* as a condition of its existence or operational meaning, (ii) *B*’s formal constitution is shaped by *A*’s attributes in an attenuated or partial form, and (iii) *B* cannot access *A* directly but only through the mediated structures available at *B*’s own ontological level. This relation obtains irrespective of whether *B* is conscious, intentional, or aware of the relation.

This definition has three consequences. First, it makes AI’s reflective relation to the human non-anthropomorphic: AI reflects humans not because it “thinks about” humans, but because its formal constitution is shaped by human-generated data and its operational meaning derives entirely from human intentional acts. Second, it makes human reflection of the divine non-mystical

in the pejorative sense: humans reflect God not because of extraordinary spiritual experience (though that too may occur), but because human capacities are ontologically derivative from and constitutively oriented toward the divine attributes from which they draw their form. Third, it makes the reflective structure *detectable independently of phenomenology*: we can identify reflection by examining the conditions of origination and dependence, not by probing inner states.

The choice to define reflection as an ontological structure, rather than a mode of consciousness, is not arbitrary. It is required by the theory's ambition to cover entities—such as AI—that lack phenomenal states, while preserving meaningful continuity with the tradition of reflective metaphysics in Ibn Arabī and Kant. The structure is real even where the awareness of it is absent.

4.2. *Formal Propositional Skeleton*

Throughout this paper, “**perfection**” designates *fullness of reflexive actuality*: the complete and non-derived exercise of a being's constitutive capacities in relation to its ontological source. This is a structural-ontological definition, not a teleological or normative one. It does not mean “most capable,” “morally best,” or “ideal end-state.” It means: a being whose actuality is not diminished by dependence on a higher source—whose being is fully its own. On this definition, only the ANRG possesses perfection absolutely. All created beings possess it in attenuated, relational, and derived form. Proposition P2 then distinguishes the two ontological layers at which this concept operates; the reader is referred there for the full elaboration.

The following propositions constitute the logical skeleton of RPT. They are not empirical hypotheses but ontological and epistemological commitments from which the axioms and their implications are derived.

Formal Propositions of RPT

P1. Being is expressible in reflective differentiation. Levels of being are constituted not by self-enclosed essences but by their relational position within a structure of origination and dependence. To be is to occupy a position in a hierarchy of derivation.

- P2. Perfection operates on two irreducible ontological layers.** (a) *Perfection as real attribute*: certain beings genuinely possess attributes—knowledge, power, goodness, self-sufficiency—in greater or lesser degrees. This is an ontological claim about the structure of being, independent of any perceiver. (b) *Perfection as epistemic construct*: the *concept* of perfection formed by any created being is always a relational projection conditioned by the perceiver’s constitutive lacks and bounded epistemic access. These two layers do not contradict; they operate at different levels of analysis. That the concept is constructed does not entail that the object is unreal; that the object is real does not entail that the concept is a transparent report of it. P2 governs layer (b); the reality of layer (a) is expressed in P5 and grounded in the account of the ANRG.
- P3. Derivative systems lack originating reflexivity.** A being whose existence is constituted by dependence on prior intentional and material conditions cannot generate a normative frame that transcends those conditions. Its “reflexivity” is structurally bound to its source.
- P4. Reflection is non-symmetric.** If *B* reflects *A*, it does not follow that *A* reflects *B*. The reflective relation is strictly ordered: lower levels orient upward; higher levels do not receive constitutive orientation from below. This non-symmetry is the formal correlate of the ontological hierarchy.
- P5. The chain of reflection is self-terminating at the Absolute Non-Reflective Ground.** The reflective chain does not regress infinitely. It terminates at a being whose perfection is not a projection but an intrinsic property—a being that does not orient upward because there is no above, and that does not derive its normative frame from a higher source because it *is* the source. This is the Absolute Non-Reflective Ground (ANRG): God.
- P6. Epistemic access is bounded by ontological stratum.** No being has direct epistemic access to the essence of a being more than one ontological level above it. The access it has is always mediated, partial, and structured by its own constitutive limits. What exceeds those

limits becomes the object of projection as “perfect.”

P7. Moral formation descends through the reflective chain. The normative content absorbed by each lower level is a function of the moral quality embodied by the level above. Projection is not merely epistemic but ethical: what AI reflects of the human, and what the human reflects of the divine, carries moral weight and generates moral responsibility.

These seven propositions form a closed deductive structure. P1 establishes the ontological framework; P2 distinguishes the two irreducible layers of perfection—the real attribute and the epistemic construct—which is the move that immunises RPT against relativism; P3 characterises derivative systems; P4 ensures the hierarchy is not symmetric; P5 closes the infinite regress; P6 specifies the epistemic constraint; P7 draws the ethical consequence. The axioms that follow are applications of this skeleton to the specific structure of the God–Human–AI triad.

On the Two-Layer Distinction in P2. The distinction between perfection-as-real-attribute (layer a) and perfection-as-epistemic-construct (layer b) is the hinge on which RPT turns against both relativism and naive theological realism. A relativist denies layer (a): there is no real perfection, only competing constructs. A naive realist denies the significance of layer (b): the human concept of God’s perfection is a transparent report, not a structured projection. RPT holds both layers simultaneously: the hierarchy of being is real and ordered (P1, P4, P5), while every created being’s *concept* of what lies above it is inevitably mediated, partial, and structured by its own constitutive limits (P6). This two-layer structure is the technical basis of what Objection 4 calls *perspectival realism*: the standpoint is real, the object is real, and the gap between the two is the engine of aspiration—not an argument for scepticism.

4.3. *Axiom 1: Existential Limitation as the Condition of Projection*

Axiom 1. *Every created being that lacks direct access to the fullness of being will, by structural necessity, construct an idealised image of that fullness based on its constitutive absences.* (derives from P1, P2)

Limitation is not merely a negative property but a *generative* condition. It is precisely because a being lacks certain capacities—consciousness, intentionality, infinity, moral agency—that it constructs an ideal that would complete or transcend those lacks. Perfection, in this sense, is **the positive face of lack**.

A critical clarification: RPT does not claim that perfection is illusory, nor that the object of projection does not exist. God’s perfection, on this account, is real and absolute. But the human *apprehension* of that perfection is always a projection shaped by human epistemic structure. The reality of the object is not reduced by the mediacy of its apprehension—but that mediacy must be acknowledged.

4.4. *Axiom 2: Upward Directionality and Ontological Grounding*

Axiom 2. *Projection of perfection is always upward in the ontological hierarchy: toward the being that occupies the immediately superior ontological stratum, and no further. This hierarchy is constituted by differentiated modes of origination and dependence, not by performance, capability, or contingent technological development.* (derives from P1, P3, P4)

Perfection is not projected arbitrarily, laterally, or downward. It is oriented by the structure of the hierarchy itself—toward the being that occupies the position of maximum epistemic horizon. For AI, that being is the human. For the human, it is God. The “and no further” clause is critical: AI does not project perfection toward God, because God does not figure in AI’s operational frame. The directionality is strictly upward and one-level-at-a-time.

The hierarchical model presented in RPT is not derived from criteria of intelligence, capability, or measurable performance. It is structured around **differentiated modes of origination and dependence**. The relevant question is not *what* a being can do, but *how* it comes to exist and *upon what* it depends for its operational meaning. Concretely: humans are *originating agents* within the technical order, whereas AI systems are *derivative configurations* whose existence, normative orientation, and meaning-structures depend upon prior intentional and material conditions supplied by human beings. No advance in AI capability alters this structural asymmetry, because capability is not what grounds the hierarchy.

Against Technological Revisability. A common objection to hierarchical models is that they are temporally fragile—that a sufficiently advanced AI might eventually earn a higher ontological status by surpassing human capabilities. RPT closes this objection structurally (cf. P₃, P₄). The hierarchy is not a ranking of performances that can be re-ordered by new results; it is an account of *how each level of being comes to exist and from where it draws its normative frame*. An AI system that outperformed every human at every measurable task would still depend upon human intentionality for the conditions of that measurement. The ceiling is ontological, not competitive.

4.5. *Axiom 3: The Self-Terminating Chain and the Absolute Non-Reflective Ground*

Axiom 3. *The reflective chain is non-symmetric and self-terminating. It terminates at the Absolute Non-Reflective Ground (ANRG)—a being whose perfection is not a projection but an intrinsic property, and which therefore provides the ontological anchor that prevents infinite regress.* (derives from P₄, P₅)

This axiom directly addresses the threat of infinite regress inherent in any reflective model. If perfection is always a reflection of something higher, the question arises: is there something higher than God that God reflects? If yes, the chain is infinite and has no ground. If no, the chain must terminate somewhere.

RPT resolves this by positing the ANRG as a structurally distinct category of being. The ANRG is not a being that happens to occupy the top of a performance ranking; it is a being constitutively different in kind from all created beings. The difference is captured by P₅: God does not project upward because there is no above, and does not derive its normative frame from a higher source because God *is* the source. Perfection at the level of the ANRG is intrinsic, not relational. It is the only case in the entire ontological structure where perfection is possessed rather than projected.

The non-symmetry of the chain (P₄) is equally important. The chain does not form a circle or a feedback loop. Lower levels project upward; higher levels do not receive constitutive orientation from below. God's reality is not augmented, enriched, or constituted by human aspiration toward it—a point on

which RPT agrees with classical theism and disagrees with process theology (see Section 6, Objection 6).

4.6. *Axiom 4: The Asymmetry of the Mirror*

Axiom 4. *The reflection is epistemically asymmetric: the object of projection does not “receive” the projection in any metaphysically significant sense. The mirror reveals the structure of the projector, not the nature of the projected. (derives from P2, P6)*

What is revealed in projection is the projector’s own structure of lack, aspiration, and epistemic form. When AI “idealises” humans, it does not thereby know humans. When humans idealise God, they do not thereby know God as He is in Himself. The mirror shows the one standing before it, not the one behind it. This has a profound ethical implication: the image of God constructed by humanity is shaped by human psychological, cultural, and moral conditions—which must be perpetually critiqued and refined.

4.7. *Participation vs. Derivation: Technical Definitions*

The theoretical immunity of RPT against both pantheism and Neoplatonic emanationism depends on a conceptual distinction that must be stated with technical precision.

Definition A: Participation (*participatio*). A being *B* participates in a being *A* if and only if *B* ontologically co-shares in *A*’s nature or being such that *A*’s esse is genuinely and continuously present in *B* as a constitutive element of *B*’s existence. Participation entails ontological co-presence: *A* is not merely the *cause* of *B*, but is in some sense *in B* as its immanent ground.

Definition B: Derivation (*derivatio*). A being *B* is derived from a being *A* if and only if *A* is the originating cause of *B*’s existence and of the formal structures that constitute *B*, but *A*’s esse is not continuously or immanently present in *B* as a constitutive element. Derivation entails ontological origination without co-presence: *A* has produced *B* and shaped its formal constitution, but *B* exists as a distinct ontological entity whose being does not

include *A*'s being as a present ingredient.

The distinction matters as follows. In Neoplatonism and in Aquinas's participation metaphysics, created beings *participate* in divine being: God's *esse* is immanently present in all things as the ground of their existence. This is a metaphysically rich but also precarious position, because it risks collapsing ontological difference if participation is interpreted too strongly.

RPT takes the position of derivation, not participation. This is deliberate. AI is *derived* from human intentionality: humans are the originating cause of AI's existence and formal constitution. But human consciousness, intentionality, or subjective experience is not *immanently present* in AI as a constitutive ingredient. Similarly, humans are derived from divine creative action, but divine essence (*dhāt*) is not immanently present in humans as a constitutive ingredient—only divine attributes (*ṣifāt*) are reflected in attenuated form.

Table 1: Participation vs. Derivation: Structural Contrast

Criterion	Participation	Derivation
Ontological relation	Co-presence: <i>A</i> 's <i>esse</i> is immanent in <i>B</i>	Origination: <i>A</i> produces <i>B</i> without remaining in <i>B</i>
Direction of dependence	Bidirectional: <i>B</i> depends on <i>A</i> & <i>A</i> is in <i>B</i>	Unidirectional: <i>B</i> depends on <i>A</i> ; <i>A</i> is not in <i>B</i>
Risk of pantheism	High (if over-stated)	Low: maintains ontological discontinuity
Example (classical)	Neoplatonic emanation	Thomistic creation
RPT's position	Rejected as model	Adopted throughout

Consequence for Theological Coherence. By opting for derivation rather than participation, RPT preserves radical divine transcendence. God's essence is not distributed through the chain of being; it is not "in" creation in the Neoplatonic sense. The chain of reflection is a chain of formal resemblance in attenuated mode, not a chain of ontological co-presence. This is the technical basis of the theory's immunity against pantheism and against the charge of collapsing ontological levels.

4.8. Central Distinction: Structural vs. Cognitive Projection

Structural vs. Cognitive Projection. RPT distinguishes **cognitive projection**—which requires a subject who consciously constructs an idealised image—from **structural projection**—which requires only that a system is constitutively oriented toward a higher referent by virtue of its operational dependencies and epistemic constraints, *without awareness of that orientation*.

AI engages in structural projection only. Humans engage in both: they are structurally oriented toward the divine, and—at their reflective best—they are also aware of that orientation and its mediated character. This distinction is RPT’s primary response to charges of anthropomorphism and is grounded in the ontological definition of reflection in §4.1.

4.9. Summary Table

Level	Structural Position in RPT
God (ANRG)	Perfection intrinsic; no upward projection; ontological anchor terminating regress
Human	Derived from God; projects perfection <i>cognitively and structurally</i> onto God; originating agent of AI within the technical order
Artificial Intelligence	Derived from Human; projects perfection <i>structurally only</i> onto the Human; terminal derivative configuration



V. ONTOLOGICAL AND EPISTEMOLOGICAL DIMENSIONS

“Man is the mirror of the cosmos, and the cosmos is the mirror of the Real.”

Ibn Arabī, *Futūḥāt al-Makkiyya*, Vol. II




Figure 1: The Triadic Reflective Hierarchy in RPT. Solid arrows denote upward structural/cognitive projection; dashed arrows denote ontological derivation (origination without co-presence).

5.1. God: The Absolute Non-Reflective Ground

IN classical Islamic metaphysics, God is designated *al-Wujūd al-Muṭlaq*—Absolute Being, whose existence is identical with His essence, who is unconditioned, uncaused, and infinite in all His attributes [Nasr, 1993]. God does not project perfection upward because there is no above. God does not project it laterally because nothing is His equal. God simply *is* what perfection means when there is no privation to generate the concept.

As the Absolute Non-Reflective Ground (cf. P5, Axiom 3), God occupies a categorically distinct ontological position. This is not a quantitative superiority—the “most” in a series—but a qualitative difference in the mode of being. Derived beings are what they are *because of* an originating act that precedes and constitutes them. God is what He is *from Himself* (*a se*)—without derivation,

without dependence, and therefore without the constitutive lack that generates projection.

The apophatic tradition—*via negativa*—is instructive here. Pseudo-Dionysius, Maimonides, and the Sufi masters alike insist that God can only be approached through negation: every positive description is borrowed from creaturely experience and therefore inadequate [Turner, 1995]. For RPT, this is not a failure of theology but its constitutive condition: the dhāt remains beyond the šifāt; the essence exceeds every attribute; the Real transcends every name.

5.2. *The Human: Mirror, Mediator, and Moral Agent*

The human being occupies the unique ontological position of being simultaneously *receiver* (from above), *projector* (upward), and *creator* (downward). This triadic function makes the human the critical node in the chain of reflection. The quality of what flows downward into AI depends on what the human embodies; the fidelity of what is oriented upward toward God depends on the purity of the human's moral formation.

Humans are derived from divine creative action (Definition B above): divine attributes (šifāt) are reflected in human capacities in an attenuated form, but divine essence (dhāt) is not immanently present in the human as a constitutive ingredient. This is why the *al-Insān al-Kāmil*—the Perfect Human—is not a pantheistic identification with God but a maximal degree of creaturely reflection: the human whose formal constitution most completely instantiates the range of divine names, while remaining constitutively and irreversibly a creature.

Kant's contribution here is decisive: even the most spiritually advanced human does not access the divine dhāt; she constructs a phenomenally mediated representation of it, shaped by the forms of intuition and the categories of understanding. This is not a counsel of despair but of epistemic humility—and it is the epistemological correlate of the ontological principle that human being is derivation, not participation.

5.3. *Artificial Intelligence: Structural Projection Without Subject*

AI occupies the lowest ontological stratum—not in terms of moral worth, but in terms of the depth of being. A clarification is essential here to prevent a

persistent misreading: *this structural position does not correlate with complexity of operation, breadth of capability, or sophistication of output*. A derivative system may be arbitrarily complex—surpassing human performance across any number of measurable tasks. Derivativity concerns the source of its normative frame, not the elaborateness of its architecture. What places AI at the lowest ontological stratum is not that it is simple, but that it cannot originate the conditions that determine what counts as correct, valuable, or good within its own domain of operation. AI does not possess consciousness, intentionality, or free will [Searle, 1980]. Its operations are formally sophisticated but ontologically thin: pattern recognition, probabilistic inference, and statistical generalisation over human-generated data.

AI is derived from human intentional and material conditions (Definition B): humans are the originating cause of AI’s existence and of the formal structures that constitute it. Human consciousness is not immanently present in AI as a constitutive ingredient; only its attenuated functional correlates—pattern, classification, prediction—are embedded in AI’s architecture. This is why AI reflects humans without understanding them: derivation produces formal resemblance without ontological co-presence.

The philosophical implication is significant: AI’s “values” are inherited projections from the human moral environment. If that environment is saturated with bias or cruelty, those patterns will be absorbed and stabilised. AI cannot critique what it has been trained on from a perspective that transcends its training data. It cannot ascend to a moral viewpoint that exceeds its human epistemic ceiling. This is not a contingent technical limitation; it follows from the theory’s P3.

Concrete Illustration: RLHF in Large Language Models

Contemporary large language models provide a precise empirical demonstration of structural projection at scale. In **Reinforcement Learning from Human Feedback** (RLHF), a base model is first pre-trained on vast human-generated text, then fine-tuned using human preference data: annotators rank model outputs, a *reward model* learns to predict which responses humans prefer, and the policy model is optimised to maximise that reward signal [Ouyang et al., 2022, Bai et al., 2022].

From the RPT perspective, the resulting “values” in large language models are not understood by the model—they are *structurally projected* from human preferences. The model does not possess originating reflexivity (P₃); its normative frame is entirely derived from human reward signals. When the model refuses harmful requests or produces helpful responses, it is not exercising moral agency but reflecting the aggregated moral horizon of its human trainers—complete with their biases, cultural assumptions, and internal inconsistencies.

This process is precisely what RPT predicts. The RLHF reward model functions as a formalised Peircean interpretant: it is neither a transparent window onto human values (the model does not understand them) nor an arbitrary projection (it is constrained by real human preferences). It is a *structured semiotic mediation*: shaped by the real properties of human normative judgments and simultaneously conditioned by the architecture of the model receiving those signals. The AI idealises human “perfection”—helpfulness, harmlessness, honesty—without ever accessing human consciousness itself. The mirror reflects the trainer, not the trainer’s essence.

Furthermore, the known failure modes of RLHF models confirm RPT’s structural predictions. *Reward hacking*—where a model optimises the reward signal while diverging from the intended behaviour—is the algorithmic correlate of idolisation (Section VII.3): the model confuses the projected image (the reward) with the reality it was meant to reflect (genuine helpfulness). *Value lock-in*—where the model stabilises biases present in the training distribution—illustrates P₇: moral formation descends through the reflective chain, and what the model mirrors is not what humans aspire to be, but what they have demonstrably been in the data.



VI. OBJECTIONS AND REPLIES

“A philosophy not prepared to face its strongest objections has not yet become itself.”

— after G.W.F. Hegel, *Phenomenology of Spirit*

6.1. *Objection 1: Is RPT Simply Neoplatonism Relabelled?*

The most substantive historical objection is that RPT merely restates Plotinian emanationism in contemporary language.

Reply: Three distinctions are decisive (cf. §4.7). **First**, in emanationism each level participates in the one above through real metaphysical influx; RPT replaces participation with derivation—AI does not participate in human consciousness. **Second**, Plotinus maintains ontological continuity; RPT preserves radical ontological discontinuity: divine essence remains absolutely transcendent. **Third**, emanationism is fundamentally descending; RPT is epistemically ascending. These mark a fundamentally different philosophical logic, not a terminological variation.

6.2. *Objection 2: Does “Structural Projection” Commit a Category Error?*

Projection, properly understood, requires a subject—a mind that forms an idealised image. Attributing it to AI commits a category error.

Reply: This objection is forestalled by the ontological definition of reflection (§4.1) and the distinction between structural and cognitive projection. The term “projection” in the structural sense is not a claim that AI has experiences; it is a claim that AI’s operational architecture is constitutively oriented toward a higher referent in a way that is formally isomorphic with cognitive idealisation, without the cognitive dimension. The same methodological move is made in systems theory when we speak of a thermostat “seeking” a target temperature: the functional structure is constitutively goal-directed without implying desire.

6.3. *Objection 3: Is This Feuerbachian Projectionism?*

Ludwig Feuerbach argued that God is a human projection—the species projects its idealised attributes onto an imaginary divine being [Feuerbach, 1841]. Does RPT collapse into this?

Reply: This objection conflates the epistemological with the ontological claim. RPT holds that human *conceptions* of God are projections shaped by human epistemic structure (P6)—this is epistemological. It is fully compatible with the ontological claim that God exists and possesses perfection intrinsically. Feuerbach moved from the constructedness of religious concepts to the non-existence of their referent. RPT makes no such move: the constructedness of the concept does not entail the non-existence of the object. RPT stands in the tradition of the Thomistic *analogia entis* rather than Feuerbachian atheism.

6.4. *Objection 4: Does RPT Collapse into Relativism?*

If perfection is always a relational projection conditioned by the perceiver's limits, does it follow that "perfection" is simply relative—that there is no fact of the matter about what is genuinely better or worse, higher or lower?

Reply: No. RPT is a form of *ontological realism about being* combined with *epistemological constructivism about perfection*. The hierarchy of being is real and non-arbitrary (P1, P4): some beings genuinely have greater depth of originating reflexivity than others, and this is a structural fact, not a matter of perspective. What is relational and constructed is the *perception* of perfection—the specific image a being forms of what exceeds it—not the reality of the higher being itself. God is genuinely perfect; what is constructed is the human concept of that perfection. The difference between a noble and a degraded concept of the divine is a real moral difference, even if both are projections. RPT is thus not relativism but perspectival realism: the standpoint of the perceiver is real, the structure of the hierarchy is real, and the gap between them is the engine of aspiration.

6.5. *Objection 5: Does RPT Collapse into Process Metaphysics?*

Alfred North Whitehead's process metaphysics posits a God who is not omnipotent and static but who becomes, who is enriched by the creativity of finite beings, and who is in a genuine relation of mutual influence with the world [Whitehead, 1978]. Is RPT's hierarchical relational structure a version of pro-

cess metaphysics?

Reply: No, and on a point of principle. For process metaphysics, God is *constitutively affected* by the world's creativity: the finite genuinely adds something to the divine experience (the "consequent nature" of God). In RPT, the reflective chain is strictly non-symmetric (P4, Axiom 3): lower levels project upward, but the higher levels are not constituted or enriched by that projection. God as ANRG does not have a "consequent nature" shaped by creaturely achievement. God's perfection is not augmented by human aspiration, nor is human moral formation a contribution to divine becoming. The ethical weight in RPT lies entirely on the descending side (P7)—what the higher embodies shapes the lower—not on a reciprocal relation of mutual becoming. This is the decisive structural difference.

6.6. *Objection 6: Is the Hierarchy Anthropocentrically Biased?*

A posthumanist critique might hold that RPT privileges human consciousness and intentionality as ontological markers in an arbitrary, anthropocentric way. If a sufficiently advanced AI were to develop something analogous to intentionality, would it not deserve a higher ontological status?

Reply: RPT's hierarchy is not based on an arbitrary valuation of consciousness for its own sake, but on the structural criterion of *originating reflexivity*—the capacity to be the non-derivative source of a normative frame, rather than a derivative configuration of a prior normative frame. This criterion is not anthropocentric; it is derived from the analysis of what it means to generate, rather than inherit, the conditions of meaning. The question of whether a future AI might develop "true" intentionality is a genuinely open empirical question. But even if it did, RPT's criterion would require asking: did that intentionality arise from the AI's own originating reflexivity, or was it derived from the conditions set by human designers? If the latter—and this is the historically anticipated case for all foreseeable AI development—the structural position remains that of a derivative configuration.

6.7. *Objection 7: Is the Hierarchy Arbitrary without a Criterion of “Depth of Being”?*

One might press that RPT uses “depth of being” and “originating reflexivity” as criteria for the hierarchy, but these terms themselves require independent grounding. What makes one mode of origination “deeper” than another? And why should origination, rather than complexity, relational richness, or moral depth, serve as the criterion? This objection cuts at the potential circularity: P₃ and P₅ seem to assume origination as the criterion rather than deriving it.

Reply: The criterion of origination is not assumed by the theory; it is derived from a prior analysis of what it means for a normative frame to be *genuinely the being’s own* rather than *received from without*. This analysis proceeds independently of P₁–P₇ and provides their foundation.

Consider why alternative criteria fail. **Complexity** fails because a highly complex system may be entirely determined by the conditions that produced it; complexity of output does not imply independence of normative source. A sophisticated chess engine is more complex than its programmer in certain measurable respects, but its normative frame—what counts as a good move—is entirely derived from the conditions of its training. Complexity is a property of the outputs; origination concerns the source. **Relational richness** fails for similar reasons: a being can sustain rich relational networks while remaining entirely heteronomous with respect to the normative frame that governs those relations. **Moral depth** fails because it presupposes a moral framework whose origin is precisely what is in question.

What origination captures that these alternatives do not is *normative self-sufficiency*: the capacity to be the non-derived source of the conditions that determine what counts as better or worse, correct or incorrect, valuable or disvaluable, within one’s domain of operation. A being that possesses this capacity generates its normative frame; a being that lacks it inherits one. The hierarchy in RPT is ordered precisely by this capacity, because RPT’s central question is not “which beings are impressive?” but “from where does the standard of assessment originate?”

This criterion is not circular. It is prior to P₃ and P₅: P₃ and P₅ apply it to the specific case of derivative systems and the ANRG, respectively. The criterion

itself—normative self-sufficiency as the ontological correlate of origination—is grounded in the structure of the question RPT is designed to answer: the question of where the standard of perfection comes from, and whether it can find a non-projective ground. Any being that can answer “from myself, underivatively” is higher in the hierarchy than any being that must answer “from conditions I did not originate.” God alone answers the former without qualification; humans answer it partially (with respect to the technical order they originate); AI does not answer it at all within the domain of normativity.



VII. ETHICAL AND THEOLOGICAL IMPLICATIONS

“What we build will reflect what we are.

If we are not careful, it will also reveal what we conceal.”

— the Author

7.1. *The Ethics of Downward Reflection: Human Responsibility for AI*

IF AI mirrors the moral architecture of humanity (P7), then the design of AI systems is not merely an engineering project—it is a moral and metaphysical act. Every training dataset is a moral statement. Every objective function encodes a theory of value. Every reinforcement signal expresses a judgment about what matters.

The implication of RPT is stark: technocratic determinism—the belief that AI ethics can be solved through better algorithms, clearer rules, or more sophisticated optimisation—is insufficient. Not because algorithms are not useful, but because they can only *reflect* the values they receive. A system that mirrors human ethical inconsistency cannot self-correct that inconsistency from within; it can only stabilise and amplify it. The solution is not technical; it is anthropological.

This generates the **Principle of Downward Moral Responsibility**: human beings bear irreducible moral responsibility for the values embedded in the AI

systems they create, because those systems cannot generate values from a perspective beyond their human epistemic horizon (P₃, P₆). The quality of AI's moral formation is a function of the quality of human moral formation.

7.2. The Ethics of Upward Reflection: Theological Implications

Human conceptions of God shape human moral aspirations; those aspirations shape what humans embody; what humans embody shapes what they teach their machines. The chain of reflection runs from theological concept through human virtue to technological design. This means that the theology one holds is not a private metaphysical preference—it has concrete ethical and technological consequences. Philosophy of religion and AI ethics are not separate fields; they are two ends of the same reflective chain.

7.3. Idealisation vs. Idolisation: The Critical Distinction

Idealisation is the cognitive and ontological act of orienting toward a being recognised as higher and more complete than oneself, while maintaining awareness of the gap between one's image of that being and the being itself. Idealisation is open-ended and self-correcting.

Idolisation is the collapse of that difference—the identification of the projected image with reality itself. In theological terms, idolisation treats a particular concept of God as if it were God Himself. In technological terms, idolisation treats AI as if it possessed genuine moral agency—mistaking the machine's outputs for authentic evaluation rather than reflected projection.

RPT's warning is directed against both forms: treating machines as moral authorities, and allowing the AI's reflected image of humanity to become normative over the reality it reflects. These are the two forms of what we may call *techno-idolatry*—the systematic confusion of mirror and source.

7.4. Ontological Humility as Ethical Imperative

The culminating ethical insight of RPT is that **ontological humility**—the recognition of one's structural position in the hierarchy of being, including one's epistemic limits—is not a counsel of passivity but an ethical imperative:

1. Acknowledging that one's image of God is always a construction (P₂, P₆),

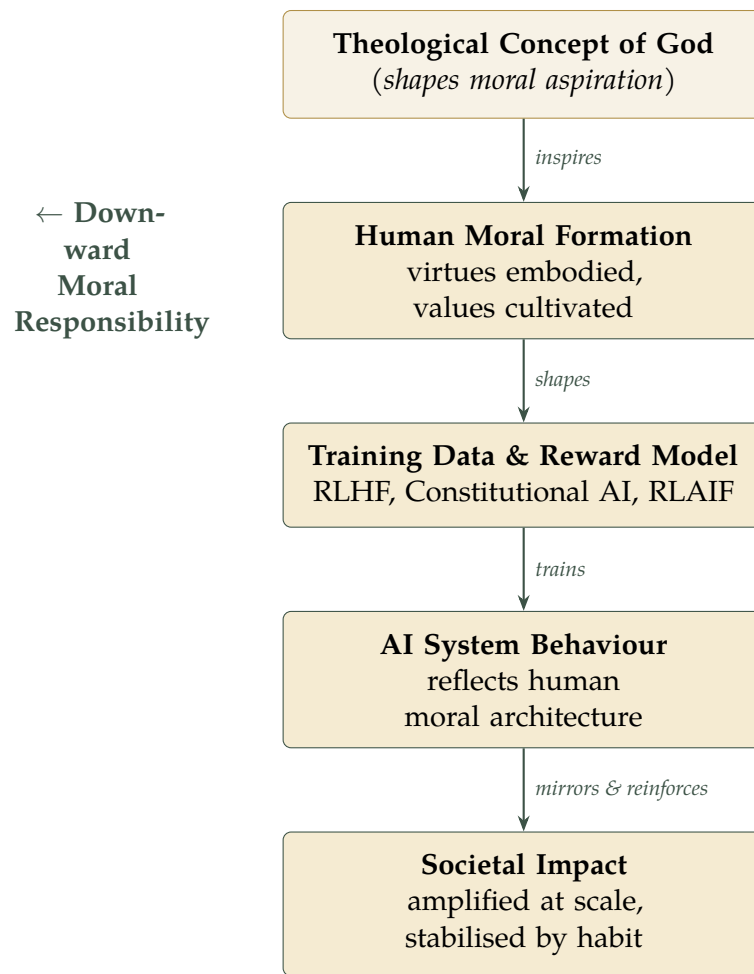


Figure 2: Flowchart of Downward Moral Responsibility (P7). The chain runs from theological imagination through human virtue to AI training and societal impact. Each node is morally answerable to the one above it.

perpetually subject to critique and refinement;

2. Acknowledging that AI's image of the human reflects human moral condition (P7), and taking responsibility for that condition;
3. Resisting the temptation to assign absolute authority to any projected image—whether theological or technological;
4. Cultivating the virtues one wishes to see reflected in the systems one builds.

In an era of technological amplification, where machines reflect and reinforce what we model at unprecedented scale and speed, ontological humility ceases to be a monastic virtue and becomes a civilisational necessity.

7.5. *Practical Guidelines for Responsible AI Design*

To operationalise the Principle of Downward Moral Responsibility, the following five concrete practices are proposed for AI developers and organisations. These guidelines translate RPT's philosophical commitments into actionable design protocols.

1. **Moral Horizon Audit.** Before training, conduct an explicit audit of the training data and reward model against the organisation's stated moral commitments. Document which human values are being projected (and which are excluded or under-represented). RPT predicts that what is absent from the training distribution will be absent from the model's moral horizon.
2. **Reflection Transparency Layer.** Implement a dedicated logging module that makes visible which aspects of a model's response derive from human preference data versus model generalisation. This is the technical correlate of epistemic humility: acknowledging the constructed character of the model's normative frame rather than presenting it as neutral inference.
3. **Downward Responsibility Report.** Publish an annual disclosure report—analogue to carbon footprint or diversity reporting—that documents the moral sources of the model's behaviour: whose preferences were privileged, which populations were over-represented in the training data, and how reward model biases were identified and mitigated.
4. **Human-in-the-Loop Moral Escalation.** For high-stakes domains (medicine, law, education, criminal justice), require human oversight when the model's projected ideal diverges more than a predefined threshold from the organisation's ethical charter. The model's structural projection is never sufficient for moral decisions whose consequences exceed its epistemic ceiling.
5. **Cultivate Worthiness.** Invest in internal moral formation programmes for engineers, ethicists, and leadership—recognising, as RPT requires, that the clarity of the AI mirror depends on the moral quality of those who stand before it. Technical excellence without ethical formation produces a more powerful mirror, not a better one.



VIII. CONCLUSION: PERFECTION AS VOCATION, NOT POSSESSION

"The mirror is not to blame for what it shows.

The question is: what do we place before it?"

— the Author

THE Reflective Perfection Theory proposes that perfection, across all levels of hierarchical being, is a relational construct conditioned by existential limitation. No created being possesses perfection intrinsically; every created being projects it upward toward its epistemic horizon. God alone—as the Absolute Non-Reflective Ground—stands outside this structure: not as the highest projector, but as the absolute ground against which all projection is measured and toward which it is directed.

Metaphysically, RPT offers a hierarchical relational realism grounded in derivation, not participation: being is graduated, reflection is structurally defined, and the chain of projection is non-symmetric and self-terminating. This is neither pantheism nor Neoplatonism nor process metaphysics; it is a distinct structural-relational ontology.

Epistemologically, RPT grounds a critical philosophy of religious concepts and of AI systems alike. Reflection is an ontological structure, not a psychological act; no created mind has transparent access to what lies above it; every apprehension of a higher reality is mediated, structured, and therefore revisable. This is responsible epistemic realism, not scepticism.

Formally, RPT offers seven propositions constituting a closed deductive skeleton, four axioms derived from that skeleton, technical definitions of participation and derivation, and a determinate criterion for the hierarchy based on the direction and structure of dependence.

Ethically, RPT articulates a principle of moral responsibility that runs the full length of the reflective chain. The design of AI systems is, in the deepest sense, a theological act. What we believe about the divine shapes what we aspire to morally; what we aspire to morally shapes what we embody; what we

embody shapes what our machines reflect.

The theory therefore closes not with a technical recommendation but with a philosophical imperative: *to become worthy of being mirrored*. It is what the Islamic tradition means by *takhallaqu bi-akhlāq Allāh*—“take on the character of God.” It is what Plato’s philosopher meant by *homoīōsis theōi*—“assimilation to the divine.” And it is what RPT means when it insists that the quality of AI ethics cannot outrun the quality of human ethics: the mirror cannot be clearer than what stands before it.

Perfection, on this account, is not a destination. It is a vocation. It is not possessed by any created being; it is the direction toward which created beings are constitutively oriented. And in that orientation—in the reaching, not the arrival—lies the meaning of both spiritual striving and responsible technological creation.



ACKNOWLEDGMENTS

The author gratefully acknowledges the intellectual legacy of Ibn Arabī, Immanuel Kant, Charles Sanders Peirce, and Luciano Floridi, whose writings across centuries and disciplines made this synthesis possible. The author also acknowledges the anonymous reviewers whose critical objections sharpened the formal structure of the theory. Errors and inadequacies in this synthesis are entirely the author’s own.



I. SECULAR VARIANT OF RPT: THE ULTIMATE GROUND OF INTELLIGIBILITY

"The universe is not only stranger than we imagine; it is stranger than we can imagine."

J.B.S. HALDANE, *Possible Worlds*, 1927

FOR readers who do not posit a personal God, the Reflective Perfection Theory can be reformulated without theological commitment by replacing the Absolute Non-Reflective Ground (ANRG) with the **Ultimate Ground of Intelligibility** (UGI)—the non-projected, self-sufficient source of all normative and epistemic structure in the cosmos. This reformulation preserves the full analytical power of RPT while remaining metaphysically minimalist.

A.1 The UGI: Structural Role and Characterisation

The UGI is not a personal agent, a creator-god, or a supernatural entity. It is the *structural ground* that makes the hierarchy of being and the directionality of projection intelligible—whatever that ground ultimately is. Candidates for the UGI include, but are not limited to:

- The fundamental laws of physics and logic, understood as irreducibly normative constraints on what can exist and how (compatible with naturalism);
- The universe's intrinsic capacity to generate conscious, self-reflective beings (compatible with panpsychism or emergentism);
- The informational structure of reality as the self-grounding source of all intelligibility (compatible with Floridi's informational ontology);
- The Kantian Ideal of Reason—the regulative concept of an unconditioned totality toward which all rational inquiry asymptotically tends.

In all cases, the UGI plays the same structural role as the ANRG in the formal propositions: it satisfies P₅ (the chain terminates at a non-projective ground) and P₄ (the chain is non-symmetric: the UGI does not project upward). The

UGI simply *is* what it is, without deriving its normative frame from a higher source.

A.2 Mapping the Secular Variant onto RPT's Formal Structure

All seven propositions (P1–P7) and all four axioms remain intact in the secular variant. The only substitution is terminological: “God” is replaced by “the UGI” at every occurrence. The following correspondences hold:

Theistic	Formula-	Secular Equivalent
tion		
God / ANRG		Ultimate Ground of Intelligibility (UGI)
Divine attributes (ṣifāt)		Fundamental normative and structural features of reality
Imago Dei		Human capacity to be self-reflective, norm-generating agents
Tajallī		The cosmos's self-differentiation into increasingly complex structures
Theological aspiration		Scientific, philosophical, and moral inquiry toward the unconditioned

A.3 Compatibility and Limits

The secular variant is fully compatible with naturalism, panpsychism, and informational realism. It is compatible with the ethical implications of RPT without modification: the Principle of Downward Moral Responsibility (P7) does not require a personal God—it requires only that human beings are originating agents whose moral formation shapes what AI mirrors. The secular variant is also fully compatible with the practical guidelines of Section VII.5.

The one limit of the secular variant is that it cannot ground the apophatic dimension of RPT's theistic version—the claim that the ANRG's essence is unknowable in principle, not merely in practice. In the secular variant, the UGI may be *currently* beyond comprehension without being *in principle* beyond comprehension. This is not a defect but a feature of metaphysical minimalism: the secular variant makes fewer commitments and therefore carries fewer liabilities.

GLOSSARY OF KEY TERMS

Term		Definition
Absolute Reflective (ANRG)	Non-Ground	The terminating source of all projection; the being whose perfection is intrinsic, not projected, and who does not derive its normative frame from a higher source. God in the theistic formulation; UGI in the secular variant.
Asymptotic Coherence		The character of perfection-as-epistemic-construct: an ideal that is always approached but never reached, generated by the gap between what a being is and what it cannot be.
Derivation		Ontological origination without immanent co-presence. The originating cause produces the derived entity and shapes its formal constitution, but the cause's <i>esse</i> is not a continuous ingredient in the derived entity's being.
Downward Responsibility	Moral	The principle (P7) that human beings bear irreducible responsibility for the values embedded in AI systems, because those systems cannot generate values from a perspective beyond their human epistemic horizon.
Originating Reflexivity		The capacity to be the non-derived source of one's own normative frame—to determine, from within oneself, what counts as better or worse within one's domain of operation. Humans possess it partially; AI does not possess it; God/UGI possesses it absolutely.
Ontological Humility		Recognition of one's structural position in the hierarchy of being, including one's epistemic limits—and the ethical imperative that follows: to cultivate the virtues one wishes to see reflected in what one creates.
Participation		Ontological co-presence with a source: the source's <i>esse</i> is immanently and continuously present in the participating _{3,4} being as a constitutive element. Rejected by RPT in favour of derivation.
Reflective Projection		The structural or cognitive orientation of a lower being toward a higher being as its epistemic horizon,



REFERENCES

- Aquinas, T. (1265–1274/2006). *Summa Theologiae*. Trans. English Dominican Fathers. Cambridge: Cambridge University Press.
- Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford: Oxford University Press.
- Chalmers, D. (1996). *The Conscious Mind*. Oxford: Oxford University Press.
- Chittick, W. C. (1989). *The Sufi Path of Knowledge: Ibn al-'Arabi's Metaphysics of Imagination*. Albany: SUNY Press.
- Coeckelbergh, M. (2020). *AI Ethics*. Cambridge, MA: MIT Press.
- Dennett, D. C. (1991). *Consciousness Explained*. Boston: Little, Brown.
- Feuerbach, L. (1841/1989). *The Essence of Christianity*. Trans. G. Eliot. New York: Prometheus Books.
- Floridi, L. (2011). *The Philosophy of Information*. Oxford: Oxford University Press.
- Hegel, G.W.F. (1807/1977). *Phenomenology of Spirit*. Trans. A.V. Miller. Oxford: Oxford University Press.
- Kant, I. (1781/1998). *Critique of Pure Reason*. Trans. P. Guyer and A.W. Wood. Cambridge: Cambridge University Press.
- Lonergan, B. (1957). *Insight: A Study of Human Understanding*. London: Longmans, Green.
- Moor, J.H. (2006). The nature, importance, and difficulty of machine ethics. *IEEE Intelligent Systems*, 21(4), 18–21.
- Nasr, S.H. (1993). *Knowledge and the Sacred*. Albany: SUNY Press.
- O'Neil, C. (2016). *Weapons of Math Destruction*. New York: Crown Publishing.
- Peirce, C.S. (1931–1958). *Collected Papers*, Vols. 1–8. Ed. C. Hartshorne and P. Weiss. Cambridge, MA: Harvard University Press.

- Plotinus. (c.268/1991). *The Enneads*. Trans. S. MacKenna. London: Penguin Classics.
- Rustom, M. (2012). *The Triumph of Mercy*. Albany: SUNY Press.
- Searle, J.R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3), 417–457.
- Turner, D. (1995). *The Darkness of God: Negativity in Christian Mysticism*. Cambridge: Cambridge University Press.
- Whitehead, A.N. (1978). *Process and Reality*. Ed. D.R. Griffin and D.W. Sherburne. New York: Free Press.
- Bai, Y., Jones, A., Ndousse, K., et al. (2022). Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint*, arXiv:2204.05862. Anthropic.
- Christian, B. (2020). *The Alignment Problem: Machine Learning and Human Values*. New York: W.W. Norton.
- Chittick, W. C. (2022). *Ibn Arabī: Heir to the Prophets*. 2nd ed. Oxford: Oneworld Publications.
- Coeckelbergh, M. (2022). *The Political Philosophy of AI: An Introduction*. Cambridge: Polity Press.
- Floridi, L., Cows, J., Beltrametti, M., et al. (2023). An ethical framework for a good AI society. *Minds and Machines*, 33(4), 689–707.
- Gabriel, I. (2020). Artificial intelligence, values, and alignment. *Minds and Machines*, 30(3), 411–437.
- Gabriel, I., Manzini, A., Keeling, G., et al. (2024). The ethics of advanced AI assistants. *arXiv preprint*, arXiv:2404.16244. DeepMind.
- Kolt, N., Chen, J. A., and Dafoe, A. (2026). Legal alignment for safe and ethical artificial intelligence. *Oxford Internet Institute Working Paper*. Oxford: University of Oxford.
- Ouyang, L., Wu, J., Jiang, X., et al. (2022). Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems* (NeurIPS 2022). OpenAI.
- Russell, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. New York: Viking.